

# Execution on Standalone and Spark Cluster

## Contents

1. Abstract
2. Introduction
3. Logistic Regression
4. Dataset details
5. HortonWorks Sandbox Creation
6. Google Cloud Platform
7. Code
8. References
9. Conclusion

### 1. Abstract:

Same program is being executed On Standalone spark context(local machine) and on Spark cluster, two programs are considered one with dataset of small size and one with dataset of large size. The program processing dataset of small size and big size executed faster on the Spark Cluster system and executed slower on Standalone system. HortonWorks platform has been used to execute the Python files.

### 2. Introduction:

Here data\_mllib\_sample\_libsvm\_data and data\_mllib\_sample\_libsvm\_data\_large datasets are considered for time consumption calculation. binomial\_logistic\_regression.py is executed with data\_mllib\_sample\_libsvm\_data dataset in Standalone mode and Cluster mode. binomial\_logistic\_regression.py is executed with data\_mllib\_sample\_libsvm\_data\_large dataset in Standalone mode and Cluster mode. Time is calculated for execution of the binomial\_logistic\_regression.py file in Standalone and Spark Cluster mode in both large and small dataset cases.

**Note: Zoom in the document to have perfect view of the snapshots being attached in the document.**

### 3. Logistic Regression:

Logistic Regression is one of the most commonly used Machine Learning algorithms. Data is being classified to class 0 or class1, binary classification using logistic regression

## Execution on Standalone and Spark Cluster

we obtain coefficients and Intercepts which are plotted to divide the data into two separate classes. Its mainly used for classification.

### **Logistic Regression Example: Spam Detection**

Spam detection is a binary classification problem where we have to classify whether the mail is spam or not. If the email is spam, we label it as 1 and if it is not spam, we label it as 0. In order to apply Logistic Regression to identify spam detection, the following features of the email are considered:

- Sender of the email
- Number of typos in the email
- Occurrence of words/phrases like "offer", "prize", "free gift", etc.

The resulting feature vector is then used to train a Logistic classifier which emits a score in the range 0 to 1. If the score is more than 0.5, we label the email as spam. Otherwise, we don't label it as spam.


### **Logistic Regression Example: Credit Card Fraud**

The Credit Card Fraud Detection is a major problem to be addressed in the banking industry because banks each year spend hundreds of millions of dollars due to fraud. When a credit card transaction happens, the bank takes a note of several factors. For instance, transaction date, transaction amount, place of transaction, type of purchase, etc. Based on these factors, a Logistic Regression model is developed to determine whether the transaction is a fraud or not. For instance, if the amount is too high and the bank knows that the concerned person never makes such high purchases, they may label it as a fraud.

Here the data is being classified, data is grouped to class 0 or Class1, Red Curve indicates the Logistic regression. Points greater than 0.5 belongs to class 1, less than 0.5 belongs to Class0, coefficients indicate the co-ordinates for plotting and Intercept indicates the intercept values at which the plot coincides with the axis.

Logistic regression is used for binary as well as categorical classification. Binary classification consists of two classes categorical classification consists of many classes.

## Execution on Standalone and Spark Cluster



Logistic regression is a popular method to predict a categorical response. It is a special case of Generalized Linear models that predicts the probability of the outcomes. In spark.ml logistic regression can be used to predict a binary outcome by using binomial logistic regression, or it can be used to predict a multiclass outcome by using multinomial logistic regression. family parameter is used to select between these two algorithms, or leave it unset and Spark will infer the correct variant.

Multinomial logistic regression is used for binary classification by setting the family param to "multinomial". It will produce two sets of **coefficients and two intercepts**.

#### 4. Dataset details:

0 128:51 129:159 130:253 131:159

1 159:124 160:253 161:255 162:63

This data is in libsvm format it is being used in the Spark framework. This a generic dataset being used in the Spark framework for classification, regression, svm problems.

Data used in the dataset is common data mainly written for the purpose of executing Machine learning algorithms. It can be bank information, Insurance information, its generic valid dataset being used to execute machine learning algorithm in Spark framework.

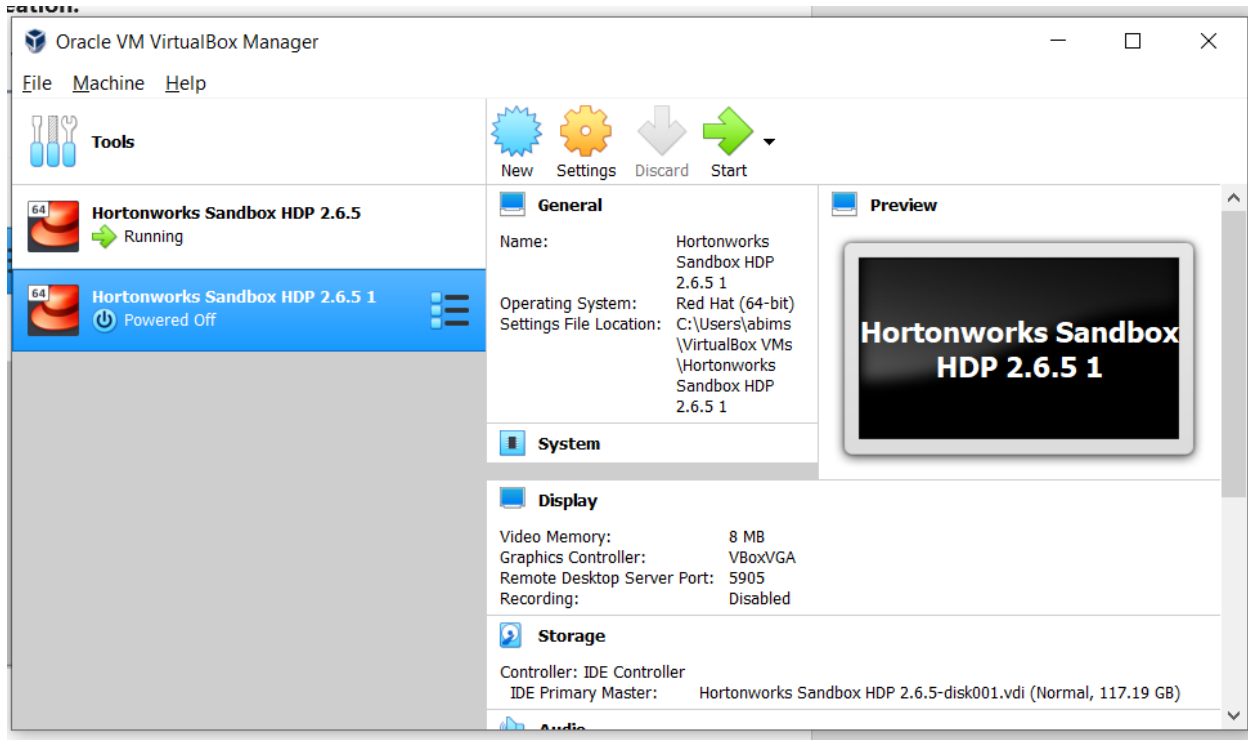
Each line contains an instance and is ended by a '\n' character. For classification, <label> is an integer indicating the class label (multi-class is supported). For regression, <label> is the target value which can be any real number. The pair <index>:<value> gives a

## Execution on Standalone and Spark Cluster

feature (attribute) value: <index> is an integer starting from 1 and <value> is a real number. The only exception is the precomputed kernel, where <index> starts from 0. Indices must be in ASCENDING order. Labels in the testing file are only used to calculate accuracy or errors. If they are unknown, just fill the first column with any numbers.

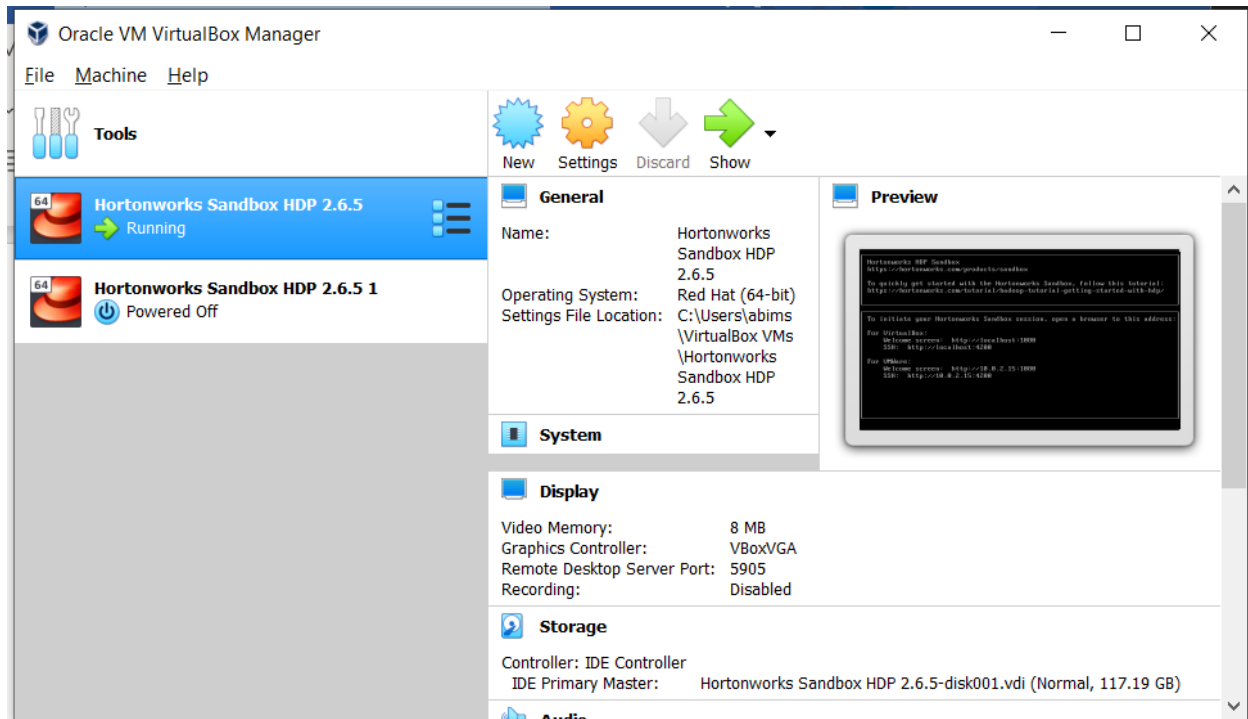
### 5. HortonsWorks Sandbox Creation:

- Creating a Hortonworks Sandbox, which provides us the platform to execute the scripts. We have to select the required Sandox and we should click on Start.



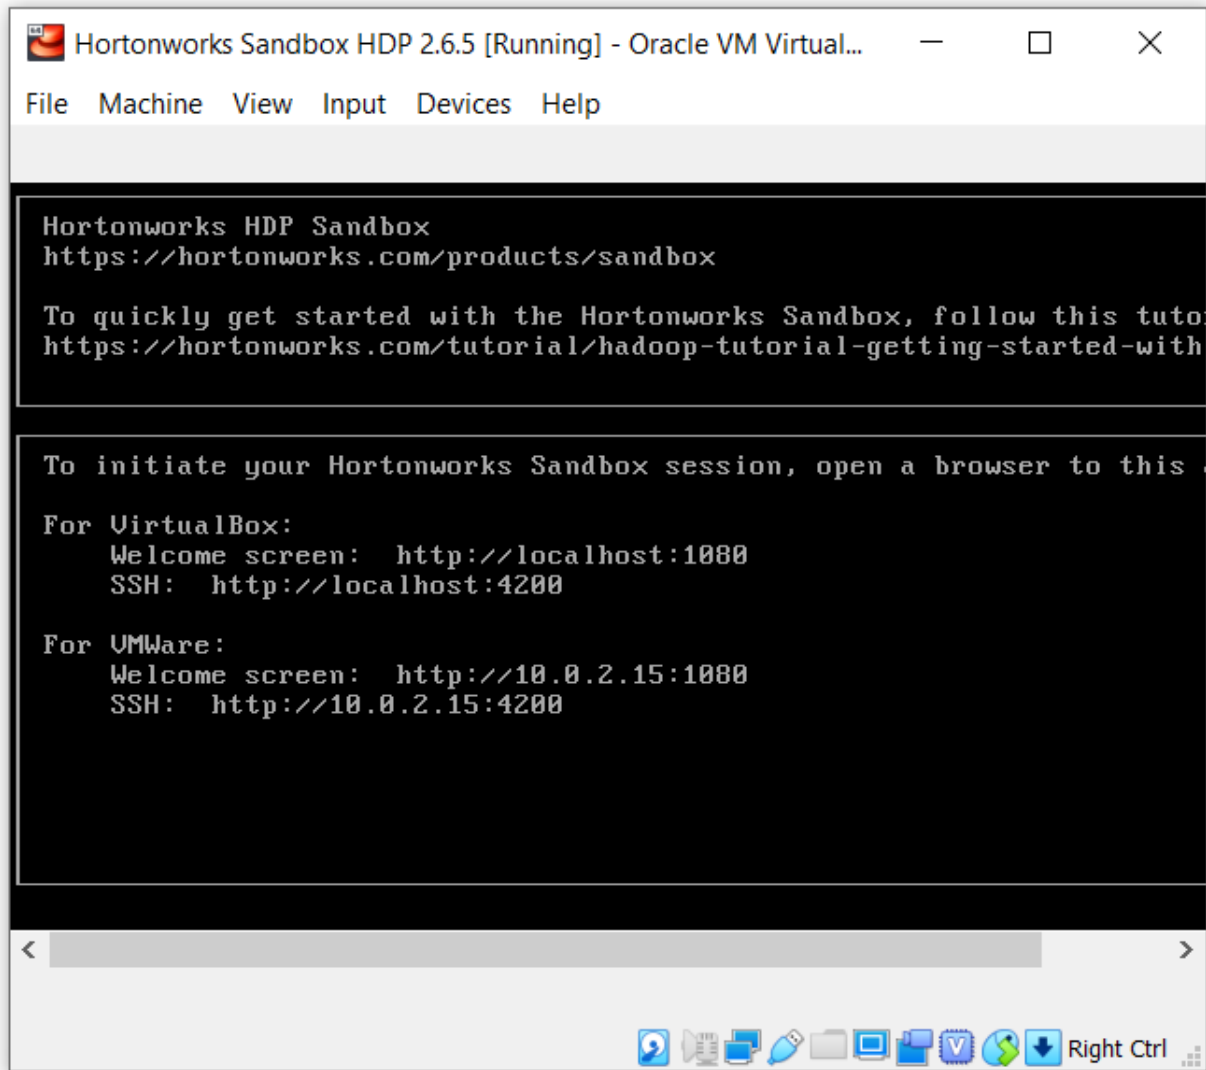
- On Starting the Sandbox VM starts running and pops up one more window with details to access the Sandbox through SSH.

## Execution on Standalone and Spark Cluster



- The below window displays Sandbox welcome screen and SSH details for VMWare and Virtualbox

## Execution on Standalone and Spark Cluster



- On Creating a Hortonworks Sandbox, We can access the terminal through SSH. We can run our scripts here.

```
sandbox-hdp login: maria_dev
maria_dev@sandbox-hdp.hortonworks.com's password:
Last login: Sun Dec 6 00:07:53 2020 from 172.18.0.2
[maria_dev@sandbox-hdp ~]$
[maria_dev@sandbox-hdp ~]$
[maria_dev@sandbox-hdp ~]$
[maria_dev@sandbox-hdp ~]$
```

- Coefficients and Intercept is required to plot Logistic regression which mainly divides the data into different classes hence Coefficients and Intercept output is printed.
- Command: **spark-submit binomial\_logistic\_regression.py**

## Execution on Standalone and Spark Cluster

Above command is executed to execute the python file with small dataset in a Standalone system.. Total time taken is in seconds not minutes.

Small Dataset:

```
20/12/06 06:03:53 INFO ContextCleaner: Cleaned accumulator 800
Multinomial coefficients: 2 X 692 CSRMatrix
(0,244) 0.0
(0,263) 0.0001
(0,272) 0.0001
(0,300) 0.0001
(0,350) -0.0
(0,351) -0.0
(0,378) -0.0
(0,379) -0.0
(0,405) -0.0
(0,406) -0.0006
(0,407) -0.0001
(0,428) 0.0001
(0,433) -0.0
(0,434) -0.0007
(0,455) 0.0001
(0,456) 0.0001
..
..
20/12/06 06:03:53 INFO BlockManagerInfo: Removed broadcast_41_piece0 on sandbox-hdp.hortonworks.com:45103 in memory (size: 8.7 KB, free: 366.1 MB)
20/12/06 06:03:53 INFO BlockManagerInfo: Removed broadcast_41_piece0 on sandbox-hdp.hortonworks.com:36145 in memory (size: 8.7 KB, free: 366.2 MB)
20/12/06 06:03:53 INFO ContextCleaner: Cleaned accumulator 820
20/12/06 06:03:53 INFO ContextCleaner: Cleaned accumulator 488
Multinomial intercepts: [-0.120658794459,0.120658794459]
('Program end : ', datetime.datetime(2020, 12, 6, 6, 3, 53, 54342))
('Total time taken in minutes', 39)
```

- Command: **spark-submit --master yarn binomial\_logistic\_regression.py**

Above command is executed to execute the python file with small dataset in a Spark Cluster system. Total time taken is in seconds not minutes.

Small Dataset:

```
Multinomial coefficients: 2 X 692 CSRMatrix
(0,244) 0.0
(0,263) 0.0001
(0,272) 0.0001
(0,300) 0.0001
(0,350) -0.0
(0,351) -0.0
(0,378) -0.0
(0,379) -0.0
(0,405) -0.0
(0,406) -0.0006
(0,407) -0.0001
(0,428) 0.0001
(0,433) -0.0
(0,434) -0.0007
(0,455) 0.0001
(0,456) 0.0001
..
..
20/12/06 06:05:09 INFO BlockManagerInfo: Removed broadcast_67_piece0 on sandbox-hdp.hortonworks.com:41675 in memory (size: 8.7 KB, free: 366.1 MB)
20/12/06 06:05:09 INFO ContextCleaner: Cleaned accumulator 537
20/12/06 06:05:09 INFO ContextCleaner: Cleaned accumulator 689
20/12/06 06:05:09 INFO ContextCleaner: Cleaned accumulator 519
20/12/06 06:05:09 INFO ContextCleaner: Cleaned accumulator 649
Multinomial intercepts: [-0.120658794459,0.120658794459]
('Program end : ', datetime.datetime(2020, 12, 6, 6, 5, 9, 299848))
('Total time taken in minutes', 6)
```

- Command: **spark-submit binomial\_logistic\_regression.py**

Above command is executed to execute the python file with small dataset in a Standalone system. Total time taken is in seconds not minutes.

Large Dataset:

# Execution on Standalone and Spark Cluster

```
20/12/06 05:55:38 INFO FileSourceScanExec: Planning scan with bin packing, max size: 81120464 bytes, open cost is considered as scanning 4194304 bytes.
20/12/06 05:55:38 INFO Instrumentation: LogisticRegression-LogisticRegression_4c46ad9ce76745d88ad3-767890594-2: training finished
Multinomial coefficients: 2 X 692 CSRMatrix
(0,244) 0.0
(0,263) 0.0001
(0,272) 0.0002
(0,300) 0.0002
(0,328) 0.0
(0,350) -0.0
(0,351) -0.0
(0,378) -0.0002
(0,379) -0.0001
(0,405) -0.0001
(0,406) -0.0005
(0,407) -0.0001
(0,428) 0.0001
(0,433) -0.0002
(0,434) -0.0005
(0,435) -0.0
..
..
Multinomial intercepts: [-0.125230028148,0.125230028148]
('Program end : ', datetime.datetime(2020, 12, 6, 5, 55, 38, 596113))
('Total time taken in minutes', 87)
20/12/06 05:55:38 INFO SparkContext: Invoking stop() from shutdown hook
```


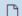
- Command: **spark-submit --master yarn binomial\_logistic\_regression.py**

Above command is executed to execute the python file with small dataset in a Standalone system. Total time taken is in seconds not minutes.

## Large Dataset:

```
20/12/06 05:51:07 INFO Instrumentation: LogisticRegression-LogisticRegression_4f8497fed357802a6507-1261717768-2: training finished
Multinomial coefficients: 2 X 692 CSRMatrix
(0,244) 0.0
(0,263) 0.0001
(0,272) 0.0002
(0,300) 0.0002
(0,328) 0.0
(0,350) -0.0
(0,351) -0.0
(0,378) -0.0002
(0,379) -0.0001
(0,405) -0.0001
(0,406) -0.0005
(0,407) -0.0001
(0,428) 0.0001
(0,433) -0.0002
(0,434) -0.0005
(0,435) -0.0
..
..
Multinomial intercepts: [-0.125230028148,0.125230028148]
('Program end : ', datetime.datetime(2020, 12, 6, 5, 51, 7, 597634))
('Total time taken in minutes', 17)
```

Dataset files uploaded to the HDFS which are being used in the Python files to perform Classification and Regression.

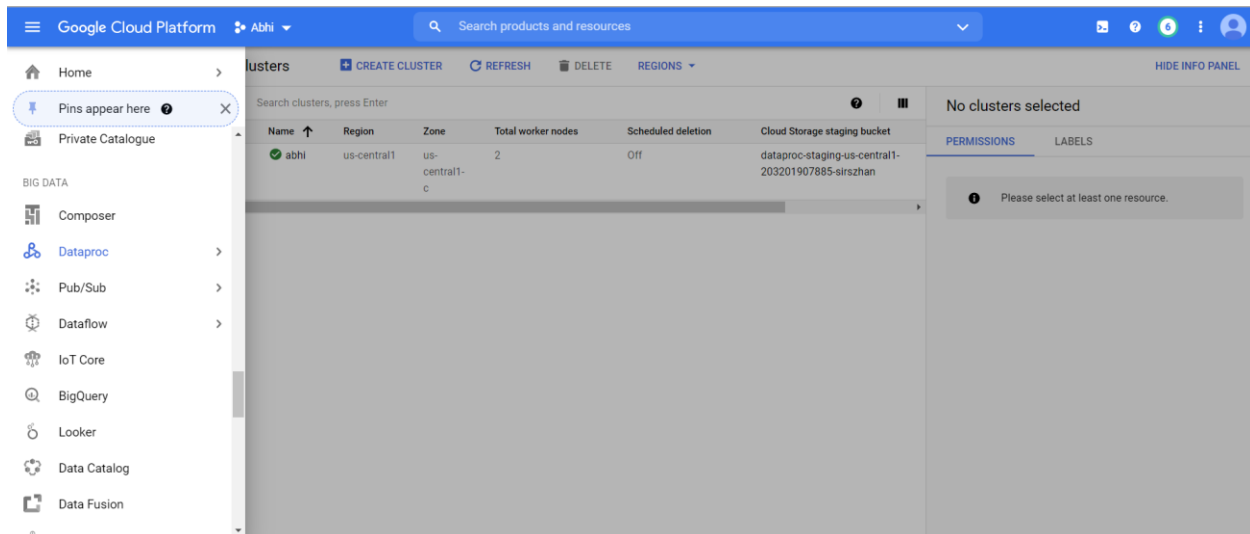
 data_mllib_sample_libsvm_data.txt	102.3 kB	2020-12-05 23:40	maria_dev	hdfs	-rw-r--r--
 data_mllib_sample_libsvm_data_large.txt	150.7 MB	2020-12-06 01:46	maria_dev	hdfs	-rw-r--r--

## 6. Google Cloud Platform :

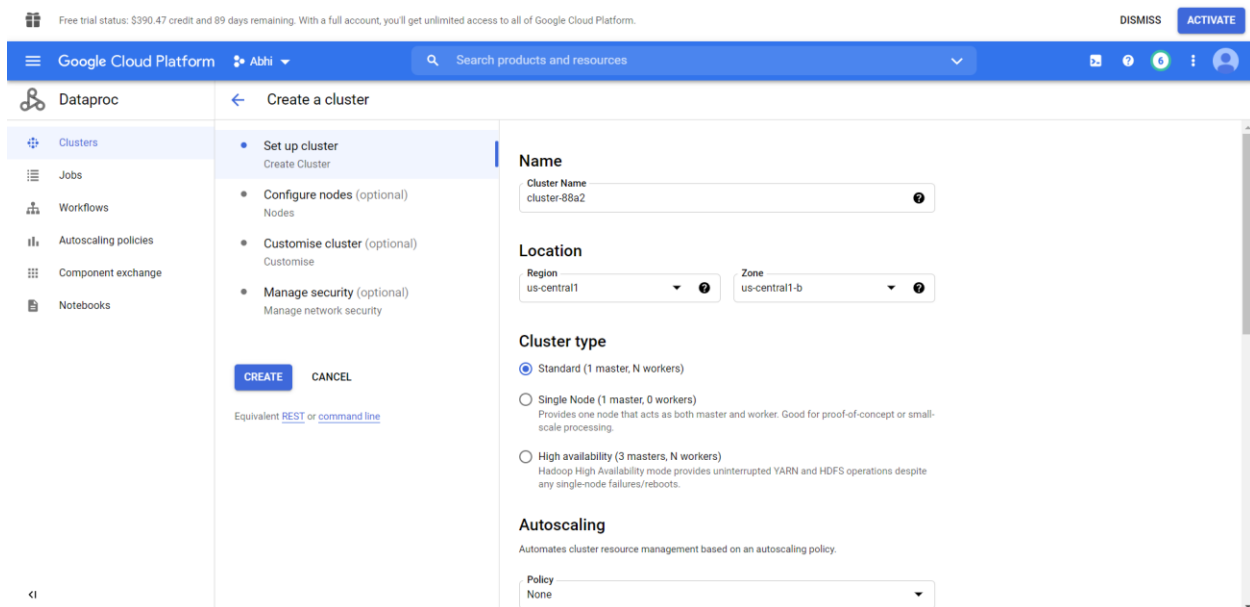
- Take the Google Cloud Platform subscription which is free for 90 days.
- Executed on Google Cloud platform as well. Dataproc is selected on the list of options present in the dropdown menu. Dataproc contains the main Cluster details. We can create our own Cluster.



# Execution on Standalone and Spark Cluster



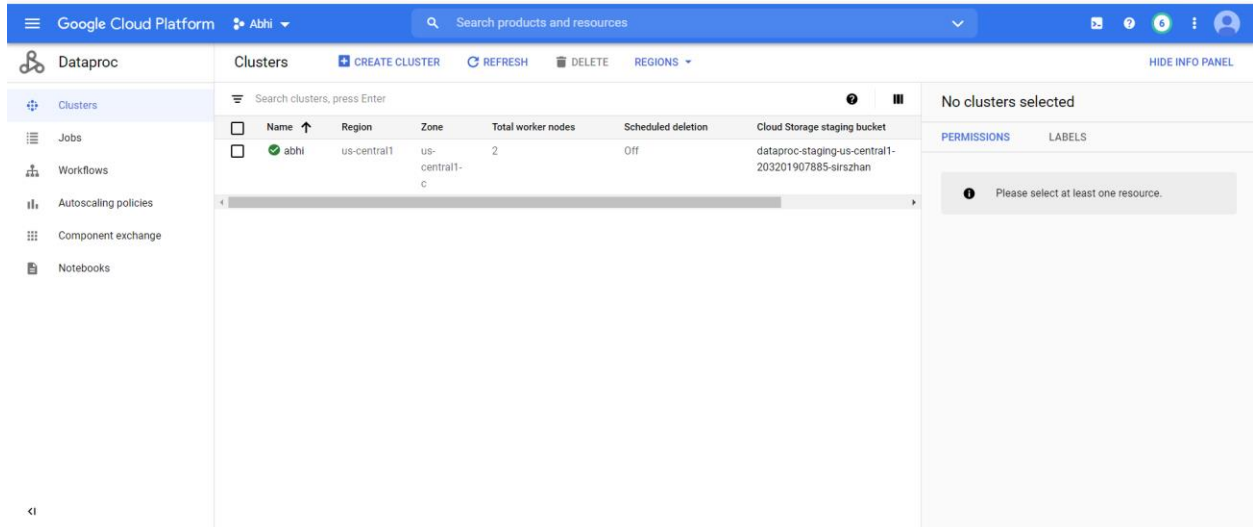
- On selecting Dataproc click on **CREATE CLUSTER** button , specify the Cluster name, Cluster type and other required details and click on create.



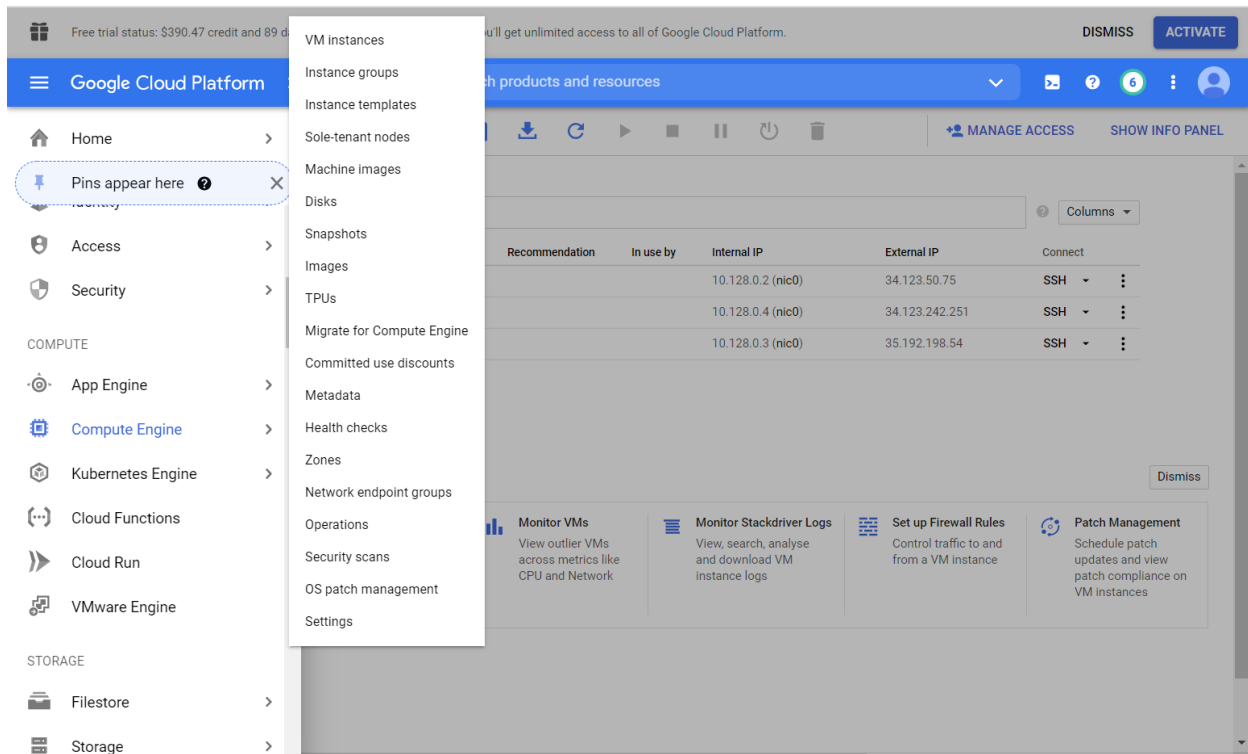
- On selecting Dataproc option we can see our new created Cluster. "Abhi" named cluster has been created here.

**Note: Zoom in the document to have perfect view of the attached snapshots in the document. Applicable to all attached snapshots.**

# Execution on Standalone and Spark Cluster

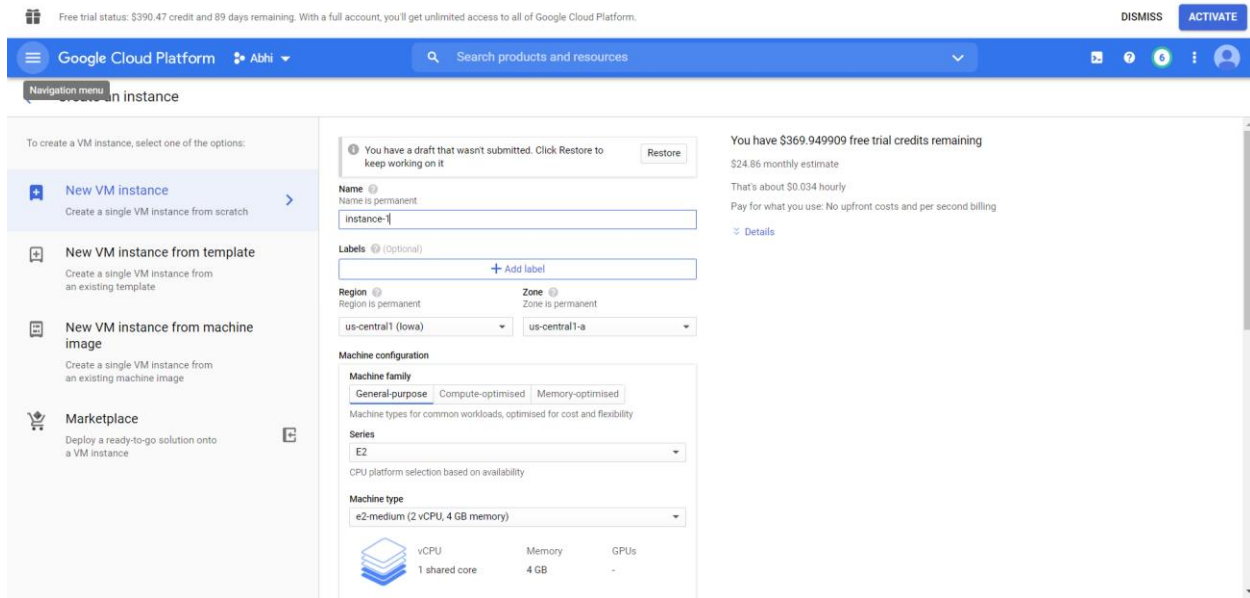


- Select "COMPUTE ENGINE" option here. It contains the main node and the worker nodes with all the details specific to master or worker nodes.

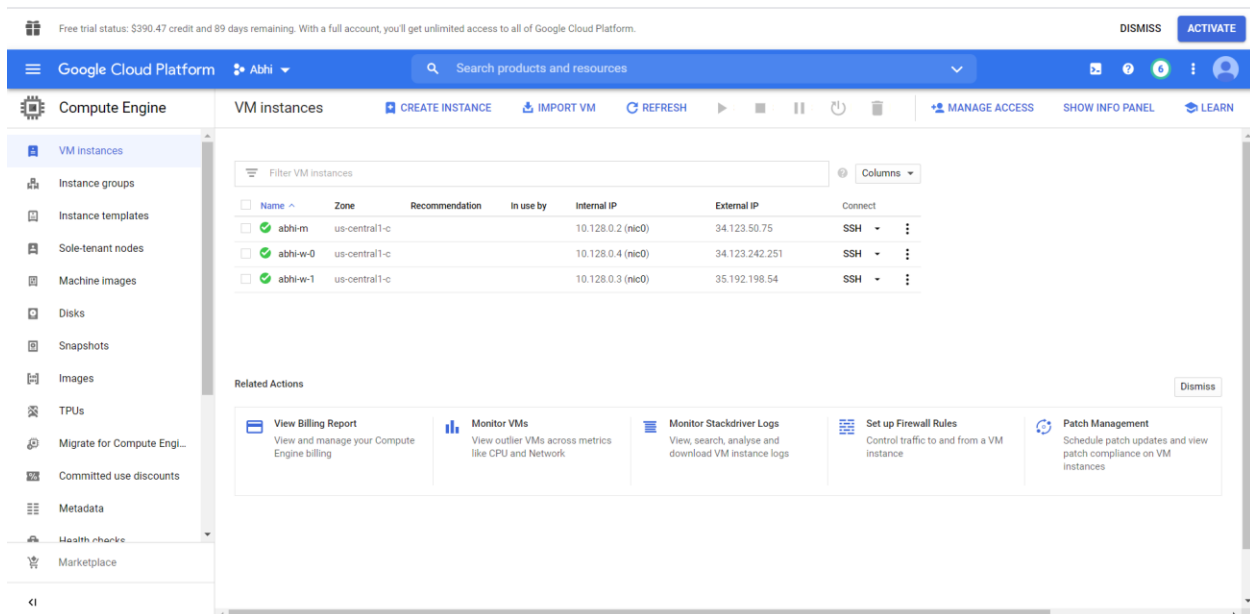


- Create VM instances with name and other required details.

# Execution on Standalone and Spark Cluster



- Select the master node "Abhi" or the worker nodes on which we are planning to execute or jobs or scripts. SSH option has been given in the picture click on SSH to access the terminal to execute scripts for master or worker nodes.



- Command: **spark-submit binomial\_logistic\_regression.py**  
Above command is executed to execute the python file with small dataset in a Standalone system.  
**Small Dataset:**

# Execution on Standalone and Spark Cluster

```
abhi@master:~$ spark-submit binomial_logistic_regression.py
(*Program started : *, datetime.datetime(2020, 12, 6, 23, 22, 27, 966993))
20/12/06 23:22:28 INFO org.spark.project.jetty.util.log: Logging initialized 82654ms
20/12/06 23:22:28 INFO org.spark.project.jetty.server.Server: jetty-9.3.z-SNAPSHOT, build timestamp: unknown, git hash: unknown
20/12/06 23:22:28 INFO org.spark.project.jetty.server.Server: Started 82748ms
20/12/06 23:22:28 INFO org.spark.project.jetty.server.AbstractConnector: Started ServerConnector876c50b5[HTTP/1.1,[http/1.1]](0.0.0.0:4040)
20/12/06 23:22:28 WARN org.apache.spark.scheduler.FairSchedulerBuilder: Fair Scheduler configuration file not found so jobs will be scheduled in FIFO order. To use fair scheduling, configure pool
is in fairscheduler.xml or set spark.scheduler.allocation.file to a file that contains the configuration.
20/12/06 23:22:30 INFO org.apache.hadoop.yarn.client.RMProxy: Connecting to ResourceManager at abhi-m/10.128.0.2:8032
20/12/06 23:22:32 INFO org.apache.hadoop.yarn.client.AHSProxy: Connecting to Application History server at abhi-m/10.128.0.2:10200
20/12/06 23:22:32 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl: Submitted application application_1607197055127_0027
20/12/06 23:22:37 WARN org.apache.spark.ml.source.libsvm.LibSVMFormat: 'numFeatures' option not specified, Determining the number of features by going though the input. If you know the number
in advance, please specify it via 'numFeatures' option to avoid the extra scan.
20/12/06 23:22:44 WARN com.github.fommil.netlib.BLAS: Failed to load implementation from: com.github.fommil.netlib.NativeSystemBLAS
20/12/06 23:22:44 WARN com.github.fommil.netlib.BLAS: Failed to load implementation from: com.github.fommil.netlib.NativeRefBLAS
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Step Size: 0.01032
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.666288 (rel: 0.0249) 1.10780
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.621707 (rel: 0.0669) 0.714252
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.612727 (rel: 0.00484) 0.349101
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.606035 (rel: 0.0109) 0.264884
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.603175 (rel: 0.00472) 0.212559
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.596962 (rel: 0.0103) 0.219945
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.594074 (rel: 0.00484) 0.192246
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.590609 (rel: 0.00583) 0.0895147
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Step Size: 0.5000
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.589472 (rel: 0.00192) 0.140102
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Step Size: 0.5000
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.588219 (rel: 0.00213) 0.0541654
20/12/06 23:22:45 INFO breeze.optimize.OWLQN: Converged because max iterations reached
Coefficients: [692, [244, 263, 272, 300, 301, 329, 350, 351, 378, 379, 405, 406, 407, 428, 433, 434, 455, 456, 461, 462, 483, 484, 489, 490, 496, 511, 512, 517, 539, 540, 568], [-7.353983524188197e-05, -9.102738505589466e-05, -0.0
0019467430546904298, -0.00020300642473486468, -3.1476183314863995e-05, -6.842977602660743e-05, 1.5883626898239883e-05, 1.4023497091372047e-05, 0.00035432047524968605, 0.0001443272898171087, 0.00010016712
070117471191634, -0.0002008459663247437, -0.0001421075579290126, 0.0002739010341160883, 0.00027730456244968115, -9.838027027269332e-05, -0.0003808522443517704, -0.00025315198008555033, 0.00027477147770754
383666666, 0.0006014109303795481, 0.0002840248179122762, -0.0001541084736508837, 0.000385996886312906, 0.000635019557424107, -0.0001506412384575676, -0.00015271865864986808, 0.0002804933808894214, 0.0006
070117471191634, -0.0002008459663247437, -0.0001421075579290126, 0.0002739010341160883, 0.00027730456244968115, -9.838027027269332e-05, -0.0003808522443517704, -0.00025315198008555033, 0.00027477147770754
307, -0.0002443619763919199, -0.0015394744687597765, -0.00023073328411331293]]
Intercept: 0.224563159613
```

```
0019467430546904298, -0.00020300642473486468, -3.1476183314863995e-05, -6.842977602660743e-05, 1.5883626898239883e-05, 1.4023497091372047e-05, 0.00035432047524968605, 0.0001443272898171087, 0.00010016712
383666666, 0.0006014109303795481, 0.0002840248179122762, -0.0001541084736508837, 0.000385996886312906, 0.000635019557424107, -0.0001506412384575676, -0.00015271865864986808, 0.0002804933808894214, 0.0006
070117471191634, -0.0002008459663247437, -0.0001421075579290126, 0.0002739010341160883, 0.00027730456244968115, -9.838027027269332e-05, -0.0003808522443517704, -0.00025315198008555033, 0.00027477147770754
307, -0.0002443619763919199, -0.0015394744687597765, -0.00023073328411331293]]
Intercept: 0.224563159613
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Step Size: 0.007313
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.660286 (rel: 0.0337) 1.43920
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.620935 (rel: 0.0596) 0.963391
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.612016 (rel: 0.0144) 0.485447
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.605265 (rel: 0.0110) 0.389485
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.602450 (rel: 0.00465) 0.291256
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Step Size: 0.5000
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.596760 (rel: 0.00944) 0.399098
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.594332 (rel: 0.00407) 0.256038
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:11:41 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.589024 (rel: 0.00893) 0.133408
20/12/06 23:11:42 INFO breeze.optimize.OWLQN: Step Size: 0.1250
20/12/06 23:11:42 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.587920 (rel: 0.00187) 0.213857
20/12/06 23:11:42 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:11:42 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.586141 (rel: 0.00302) 0.168430
20/12/06 23:11:42 INFO breeze.optimize.OWLQN: Converged because max iterations reached
Multinomial coefficients: 2 X 692 CSRMatrix
(0,244) 0.0
(0,263) 0.0001
(0,272) 0.0001
(0,300) 0.0001
(0,350) -0.0
(0,351) -0.0
(0,378) -0.0
(0,379) -0.0
(0,405) -0.0
(0,406) -0.0006
(0,407) -0.0001
(0,428) 0.0001
(0,433) -0.0
(0,434) -0.0007
(0,455) 0.0001
(0,456) 0.0001
..
..
Multinomial intercepts: [-0.12065879445860686, 0.12065879445860686]
(*Program end : *, datetime.datetime(2020, 12, 6, 23, 11, 42, 261514))
(*Total time taken in seconds: 70)
20/12/06 23:11:42 INFO org.spark.project.jetty.server.AbstractConnector: Stopped SParke18563e84[HTTP/1.1,[http/1.1]](0.0.0.0:4040)
```

- Command: **spark-submit --deploy-mode cluster binomial\_logistic\_regression.py**

Above command is executed to execute the python file with small dataset in a Spark Cluster system.

## Small Dataset:

```
abhi@master:~$ spark-submit --master yarn --deploy-mode cluster binomial_logistic_regression.py
20/12/06 23:23:50 INFO org.apache.hadoop.yarn.client.RMProxy: Connecting to ResourceManager at abhi-m/10.128.0.2:8032
20/12/06 23:23:50 INFO org.apache.hadoop.yarn.client.AHSProxy: Connecting to Application History server at abhi-m/10.128.0.2:10200
20/12/06 23:23:52 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl: Submitted application application_1607197055127_0028
abhi@master:~$
```

- Command: **spark-submit binomial\_logistic\_regression.py**

Above command is executed to execute the python file with small dataset in a Standalone system.

## Large Dataset:

# Execution on Standalone and Spark Cluster

```
abhi@testdy@abhi:~$ spark-submit binomial_logistic_regression.py
(*Program started : *, datetime.datetime(2020, 12, 6, 23, 27, 26, 570818))
20/12/06 23:27:29 INFO org.spark_project.jetty.util.log: Logging initialized @2581ms
20/12/06 23:27:29 INFO org.spark_project.jetty.server.Server: jetty-9.3.2-SNAPSHOT, build timestamp: unknown, git hash: unknown
20/12/06 23:27:29 INFO org.spark_project.jetty.server.Server: Started @2671ms
20/12/06 23:27:29 INFO org.spark_project.jetty.server.AbstractConnector: Started ServerConnector@10238cb6[HTTP/1.1,[http/1.1]](0.0.0.0:4040)
20/12/06 23:27:29 WARN org.apache.spark.scheduler.FairSchedulerBuilder: Fair Scheduler configuration file not found so jobs will be scheduled in FIFO order. To use fair scheduling, configure poo
ls in fairScheduler.xml or set spark.scheduler.allocation.file to a file that contains the configuration.
20/12/06 23:27:30 INFO org.apache.hadoop.yarn.client.RMProxy: Connecting to ResourceManager at abhi-m/10.128.0.2:8032
20/12/06 23:27:30 INFO org.apache.hadoop.yarn.client.AHSProxy: Connecting to Application History server at abhi-m/10.128.0.2:10200
20/12/06 23:27:32 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl: Submitted application application_1607197055127_0029
20/12/06 23:27:39 WARN org.apache.spark.ml.source.libsvm.LibSVMWrapper: 'numFeatures' option not specified, determining the number of features by going through the input. If you know the number
in advance, please specify it via 'numFeatures' option to avoid the extra scan.
20/12/06 23:27:54 WARN com.github.fommil.netlib.BLAS: Failed to load implementation from: com.github.fommil.netlib.NativeSystemBLAS
20/12/06 23:27:54 WARN com.github.fommil.netlib.BLAS: Failed to load implementation from: com.github.fommil.netlib.NativeRefBLAS
20/12/06 23:27:54 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.665895 (rel: 0.0255) 1.12134
20/12/06 23:27:54 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:27:54 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.620870 (rel: 0.0676) 0.714166
20/12/06 23:27:54 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:27:54 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.611804 (rel: 0.0146) 0.348395
20/12/06 23:27:55 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:27:55 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.605053 (rel: 0.0110) 0.264093
20/12/06 23:27:55 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:27:55 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.602072 (rel: 0.00493) 0.212317
20/12/06 23:27:55 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:27:55 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.595825 (rel: 0.0104) 0.223480
20/12/06 23:27:55 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:27:55 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.592955 (rel: 0.00482) 0.189788
20/12/06 23:27:55 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:27:55 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.589426 (rel: 0.00595) 0.0907327
20/12/06 23:27:56 INFO breeze.optimize.OWLQN: Step Size: 0.5000
20/12/06 23:27:56 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.588378 (rel: 0.00178) 0.144383
20/12/06 23:27:56 INFO breeze.optimize.OWLQN: Step Size: 0.5000
20/12/06 23:27:56 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.587166 (rel: 0.00206) 0.0580350
20/12/06 23:27:56 INFO breeze.optimize.OWLQN: Converged because max iterations reached
Coefficients: [692,1244,263,272,300,301,328,350,351,378,379,405,406,407,428,433,434,455,456,461,462,483,484,489,490,496,511,512,517,539,540,560],[7.31158629541713e-05,-9.72461824132869e-05,-0.0
019622622352728143,-0.00020452700113987916,-3.775920041728073e-05,-7.193406184967735e-05,1.472519739214789e-05,1.360980279403029e-05,0.0003464164431832787,0.0001159025465737444,0.000100966520248
3371,0.0006133762812073982,0.0002849504490361429,-0.00012194405580543064,0.00038170519747171506,0.0006476659835453523,-0.00012175992138273959,-0.0001595332713816548,0.00028113293802198546,0.000619
2370823118477,-0.0002034800753947824,-0.00015132056330607738,0.0002746340937965658,0.0002782934171066985,-0.0001042560711674442,-0.00035116848044106167,-0.0002046859651260234,0.00027841868802051
204,-0.00023206871716881082,-0.001496671156386187,-0.00021881653385004928]]
Intercept: 0.22417691612
+
ssh.cloud.google.com/projects/abhi-297719/zones/us-central1-c/instances/abhi-m?useAdminProxy=true&authuser=1&hl=en_GB&projectNumber=203201907885
0019622622352728143,-0.00020452700113987916,-3.775920041728073e-05,-7.193406184967735e-05,1.472519739214789e-05,1.360980279403029e-05,0.0003464164431832787,0.0001159025465737444,0.000100966520248
3371,0.0006133762812073982,0.0002849504490361429,-0.00012194405580543064,0.00038170519747171506,0.0006476659835453523,-0.00012175992138273959,-0.0001595332713816548,0.00028113293802198546,0.000619
2370823118477,-0.0002034800753947824,-0.00015132056330607738,0.0002746340937965658,0.0002782934171066985,-0.0001042560711674442,-0.00035116848044106167,-0.0002046859651260234,0.00027841868802051
204,-0.00023206871716881082,-0.001496671156386187,-0.00021881653385004928]]
Intercept: 0.22417691612
+
20/12/06 23:28:01 INFO breeze.optimize.OWLQN: Step Size: 0.007264
20/12/06 23:28:01 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.659746 (rel: 0.0345) 1.45519
20/12/06 23:28:01 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:28:01 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.620379 (rel: 0.0597) 0.963758
20/12/06 23:28:01 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:28:01 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.611376 (rel: 0.0145) 0.489569
20/12/06 23:28:01 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:28:01 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.604455 (rel: 0.0113) 0.388002
20/12/06 23:28:01 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:28:01 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.601517 (rel: 0.00486) 0.289814
20/12/06 23:28:02 INFO breeze.optimize.OWLQN: Step Size: 0.5000
20/12/06 23:28:02 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.595949 (rel: 0.00926) 0.421431
20/12/06 23:28:02 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:28:02 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.593665 (rel: 0.00383) 0.230670
20/12/06 23:28:02 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:28:02 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.590195 (rel: 0.00584) 0.243256
20/12/06 23:28:02 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:28:02 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.589410 (rel: 0.00133) 0.325256
20/12/06 23:28:02 INFO breeze.optimize.OWLQN: Step Size: 1.000
20/12/06 23:28:02 INFO breeze.optimize.OWLQN: Val and Grad Norm: 0.587059 (rel: 0.00399) 0.230659
20/12/06 23:28:02 INFO breeze.optimize.OWLQN: Converged because max iterations reached
Multinomial coefficients: 2 X 632 CoefMatrix
(0.244) 0.0
(0.263) 0.0001
(0.272) 0.0002
(0.300) 0.0002
(0.328) 0.0
(0.350) -0.0
(0.351) -0.0
(0.378) -0.0002
(0.379) -0.0001
(0.405) -0.0001
(0.406) -0.0005
(0.407) -0.0001
(0.428) 0.0001
(0.433) -0.0002
(0.434) -0.0005
(0.435) -0.0
..
Multinomial intercepts: [-0.12523002014839235,0.12523002014839235]
(*Program end : *, datetime.datetime(2020, 12, 6, 23, 28, 3, 93604))
(*Total time taken in seconds*: 34)
20/12/06 23:28:03 INFO org.spark_project.jetty.server.AbstractConnector: Stopped Spark@10238cb6[HTTP/1.1,[http/1.1]](0.0.0.0:4040)
abhi@testdy@abhi:~$
```

- Command: **spark-submit --deploy-mode cluster binomial\_logistic\_regression.py**

Above command is executed to execute the python file with small dataset in a Spark Cluster system.


## Large Dataset:

```
abhi@testdy@abhi:~$ spark-submit --master yarn --deploy-mode cluster binomial_logistic_regression.py
20/12/07 00:22:17 INFO org.apache.hadoop.yarn.client.RMProxy: Connecting to ResourceManager at abhi-m/10.128.0.2:8032
20/12/07 00:22:18 INFO org.apache.hadoop.yarn.client.AHSProxy: Connecting to Application History server at abhi-m/10.128.0.2:10200
20/12/07 00:22:20 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl: Submitted application application_1607197055127_0030
abhi@testdy@abhi:~$
```

<http://34.123.50.75:8088/cluster> is clicked to access the below attached screen which shows the execution details of the python script being executed and all other required cluster details.

# Execution on Standalone and Spark Cluster

← → ↻ ⚠ Not secure | 34.123.50.75:8088/cluster



## All Applications

Cluster

About

Nodes

Node Labels

Applications

NEW

NEW SAVING

SUBMITTED

ACCEPTED

RUNNING

FINISHED

FAILED

KILLED

Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Re
29	0	0	29	0	0 B	24 GB	0 B

Cluster Nodes Metrics

Active Nodes	Decommissioning Nodes	Decommissioned Nodes	Lost Nodes	Unhealthy Nodes
2	0	0	0	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[MEMORY]	<memory:1024, vCores:1>	<memory:12288, vCores:4>

Show 20 ▾ entries

ID	User	Name	Application Type	Queue	Application Priority	StartTime	FinishTime	State	FinalStatus	Running Containers	Allocated CPU V-Cores	All...
application_1607197055127_0029	abimsstudy	ML-DEMO	SPARK	default	0	Sun Dec 6 19:27:32 -0400 2020	Sun Dec 6 19:28:03 -0400 2020	FINISHED	SUCCEEDED	N/A	N/A	N/A
application_1607197055127_0028	abimsstudy	binomial_logistic_regression.py	SPARK	default	0	Sun Dec 6 19:23:52 -0400 2020	Sun Dec 6 19:24:16 -0400 2020	FINISHED	SUCCEEDED	N/A	N/A	N/A
application_1607197055127_0027	abimsstudy	ML-DEMO	SPARK	default	0	Sun Dec 6 19:22:32 -0400 2020	Sun Dec 6 19:22:51 -0400 2020	FINISHED	SUCCEEDED	N/A	N/A	N/A
application_1607197055127_0026	abimsstudy	ML-DEMO	SPARK	default	0	Sun Dec 6 19:11:26 -0400 2020	Sun Dec 6 19:11:42 -0400 2020	FINISHED	SUCCEEDED	N/A	N/A	N/A
application_1607197055127_0027	abimsstudy	ML-DEMO	SPARK	default	0	Sun Dec 6 19:23:32 -0400 2020	Sun Dec 6 19:22:51 -0400 2020	FINISHED	SUCCEEDED	N/A	N/A	N/A
application_1607197055127_0026	abimsstudy	ML-DEMO	SPARK	default	0	Sun Dec 6 19:11:26 -0400 2020	Sun Dec 6 19:11:42 -0400 2020	FINISHED	SUCCEEDED	N/A	N/A	N/A
application_1607197055127_0025	abimsstudy	binomial_logistic_regression.py	SPARK	default	0	Sun Dec 6 19:09:47 -0400 2020	Sun Dec 6 19:10:24 -0400 2020	FINISHED	SUCCEEDED	N/A	N/A	N/A
application_1607197055127_0024	abimsstudy	ML-DEMO	SPARK	default	0	Sun Dec 6 17:55:53 -0400 2020	Sun Dec 6 17:56:20 -0400 2020	FINISHED	SUCCEEDED	N/A	N/A	N/A
application_1607197055127_0023	abimsstudy	ML-DEMO	SPARK	default	0	Sun Dec 6 17:55:09 -0400 2020	Sun Dec 6 17:55:36 -0400 2020	FINISHED	SUCCEEDED	N/A	N/A	N/A
application_1607197055127_0022	abimsstudy	binomial_logistic_regression.py	SPARK	default	0	Sun Dec 6 17:54:09 -0400 2020	Sun Dec 6 17:54:47 -0400 2020	FINISHED	SUCCEEDED	N/A	N/A	N/A
application_1607197055127_0021	abimsstudy	ML-DEMO	SPARK	default	0	Sun Dec 6 17:26:04 -0400 2020	Sun Dec 6 17:26:34 -0400 2020	FINISHED	SUCCEEDED	N/A	N/A	N/A
application_1607197055127_0020	abimsstudy	binomial_logistic_regression.py	SPARK	default	0	Sun Dec 6 17:01:46 -0400 2020	Sun Dec 6 17:02:20 -0400 2020	FINISHED	SUCCEEDED	N/A	N/A	N/A

## 7. Code:

```
from pyspark.ml.classification import LogisticRegression

from datetime import datetime

start_time = datetime.now()

print("Program started : ",start_time)
```

## Execution on Standalone and Spark Cluster

```
spark = SparkSession.builder.appName("ML-DEMO").getOrCreate()

#training_lr = spark.read.format("libsvm").load("data_mllib_sample_libsvm_data.txt")

# Load training data

training_lr = spark.read.format("libsvm").load("data_mllib_sample_libsvm_data_large.txt")

logreg = LogisticRegression(maxIter=10, regParam=0.3, elasticNetParam=0.8)

# Fit the model

lrModel = logreg.fit(training_lr)

# Print the coefficients and intercept for logistic regression

print("Coefficients: " + str(lrModel.coefficients))

print("Intercept: " + str(lrModel.intercept))

# We can also use the multinomial family for binary classification

mlogreg = LogisticRegression(maxIter=10, regParam=0.3, elasticNetParam=0.8,
family="multinomial")

# Fit the model

mlrModel = mlogreg.fit(training_lr)

# Print the coefficients and intercepts for logistic regression with multinomial family

print("Multinomial coefficients: " + str(mlrModel.coefficientMatrix))

print("Multinomial intercepts: " + str(mlrModel.interceptVector))

end_time = datetime.now()

print("Program end : ",end_time)

time_taken = end_time -start_time

print("Total time taken in minutes",int(time_taken.seconds))
```

### 8. References:

[https://scikit-learn.org/stable/auto\\_examples/linear\\_model/plot\\_logistic.html](https://scikit-learn.org/stable/auto_examples/linear_model/plot_logistic.html)

<https://spark.apache.org/docs/latest/mllib-linear-methods.html#logistic-regression>

### 9. Conclusion:

Here binomial\_logistic\_regression.py is executed with small and large datasets when its executes on Standalone system more time is consumed and when we run on Spark cluster where execution is divided among clusters and less time is consumed.